

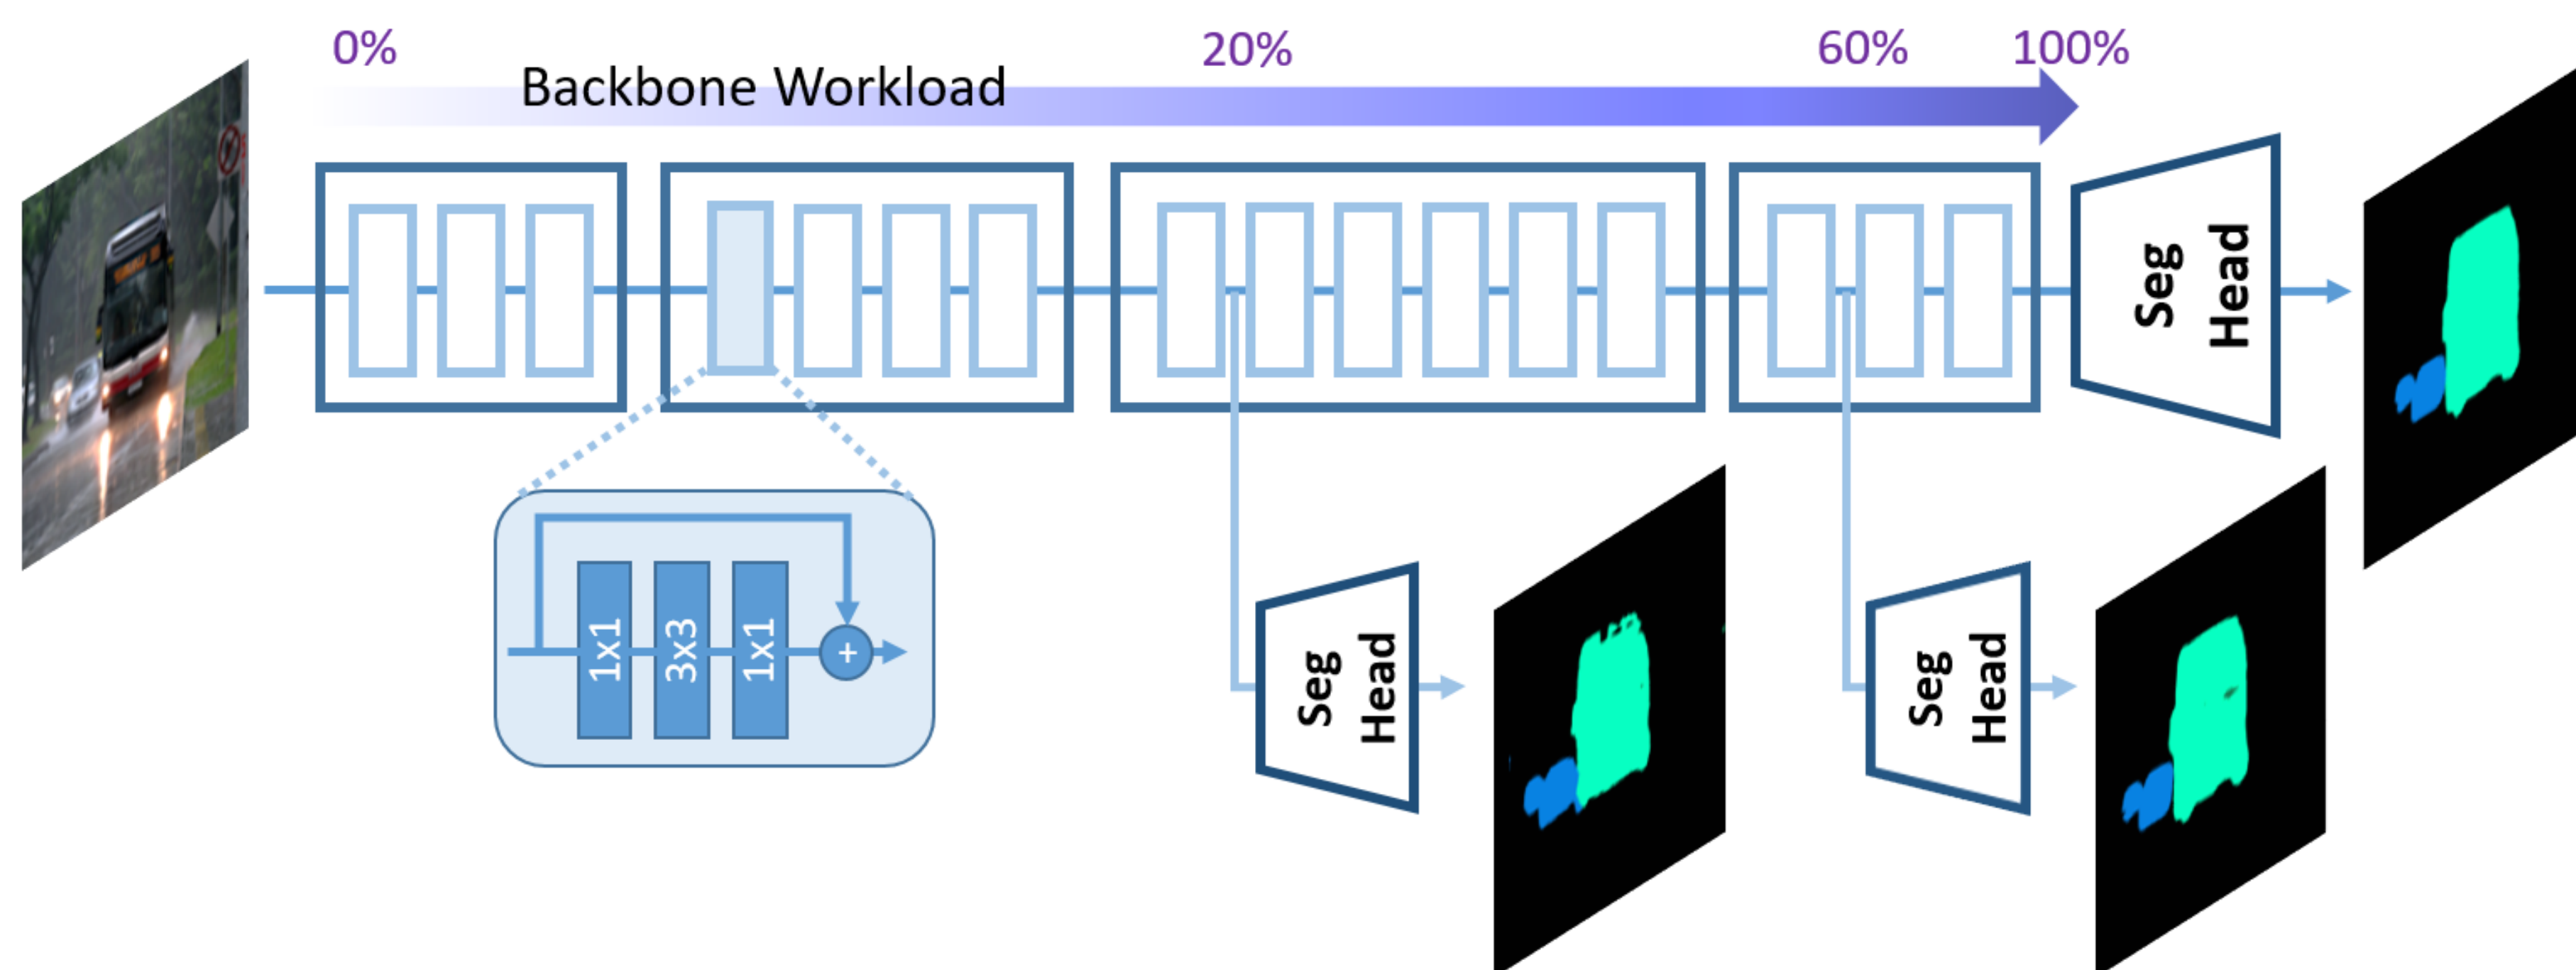
Overview

MESS is a [framework](#) that converts segmentation CNNs to multi-exit models:

- Specially trained models that employ parametrised exits along their depth.
- This allows to dynamically save computation during inference on easier samples.
- MESS jointly optimises the number, placement and architecture of early-exits.
- Saves training and maintenance cost by allowing post-training adaptation.

To enable this we also introduce:

- A [novel two-stage training](#) scheme for MESS networks, combining exit-aware backbone pre-training, with a selective self-distillation loss for the exits.
- An input-dependent inference pipeline for MESS networks, introducing a novel [exit-policy](#), tailored to dense-output models, such as semantic segmentation.



Key findings

Compared to their respective backbones MESS networks are able to achieve:

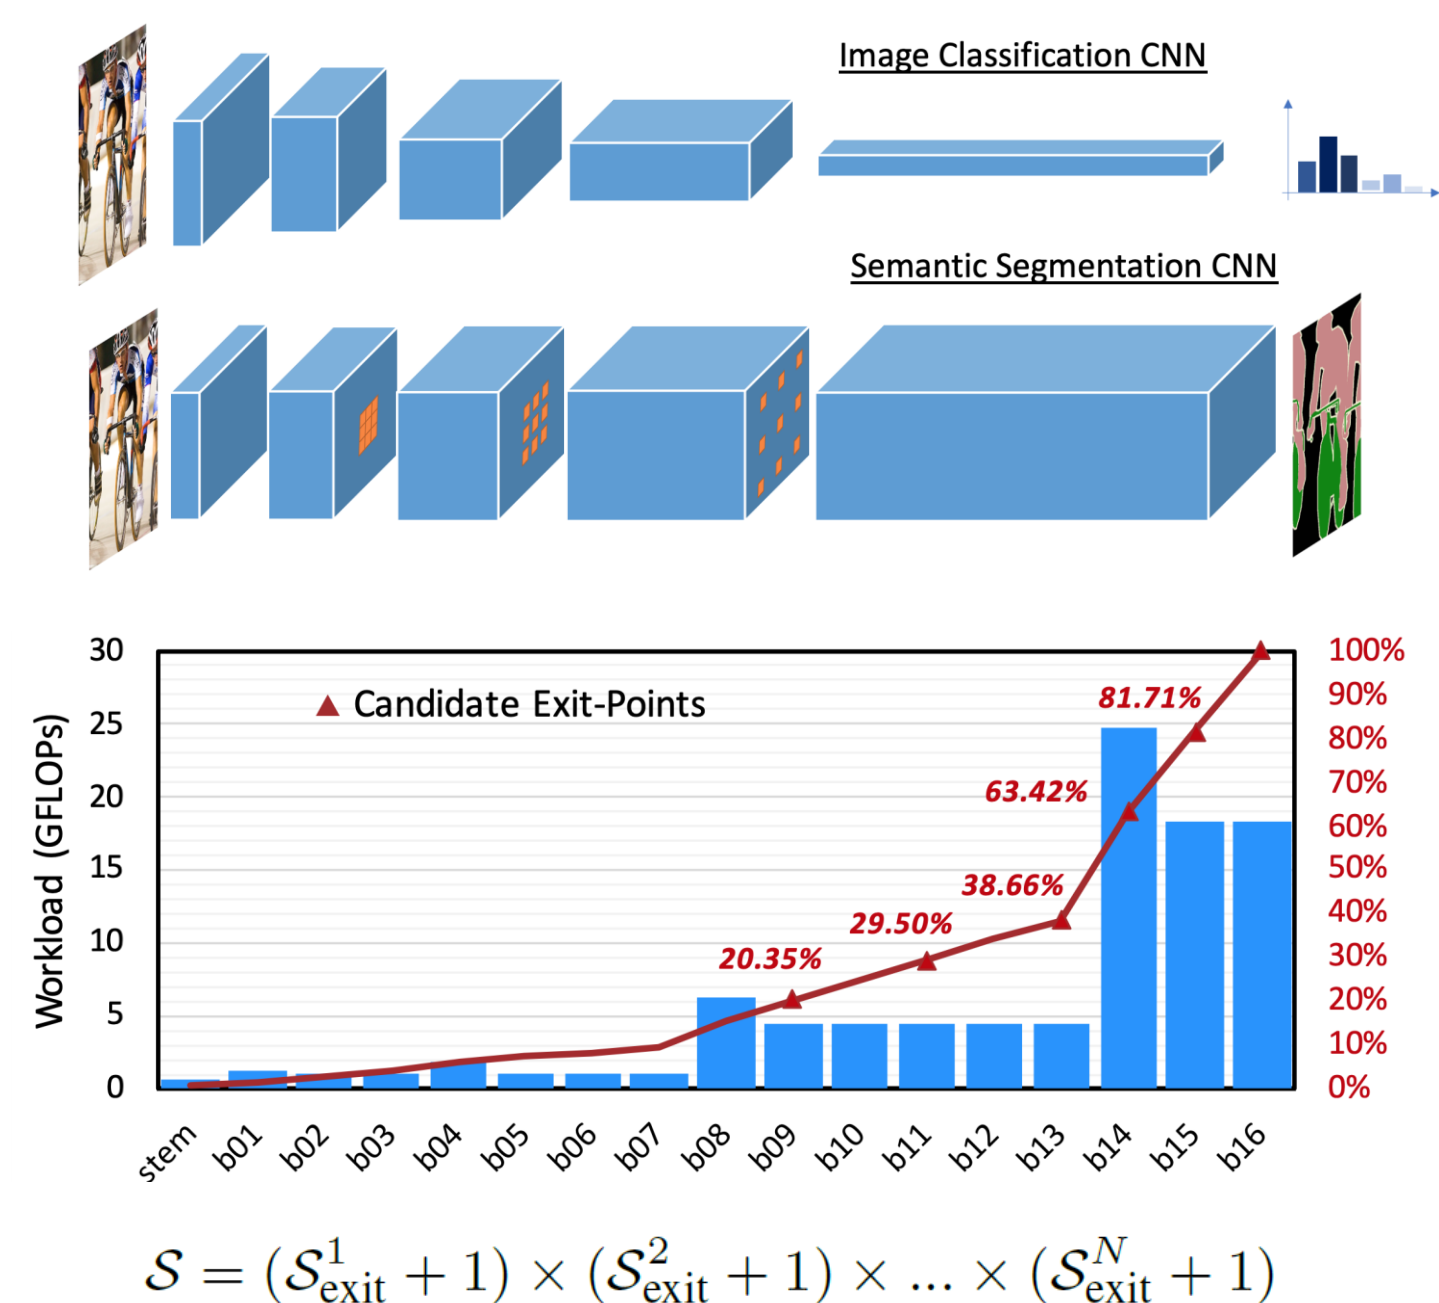
- Up to **2.8x faster inference**, with the same accuracy
- Up to **5.3 pp higher accuracy**, for the same computation budget

MESS networks also offer:

- Better (but comparable) speed-accuracy trade-off than SOTA NAS-crafted models, while being **10x faster to train**.
- Up to **3x faster inference** than uniform-exit MESS instances (status-quo).
- Gains that translate to **realistic latency speed-ups** out-of-the-box.
- Remain **adaptable post-training** via search (<1GPUh).

Candidate Exit-Points

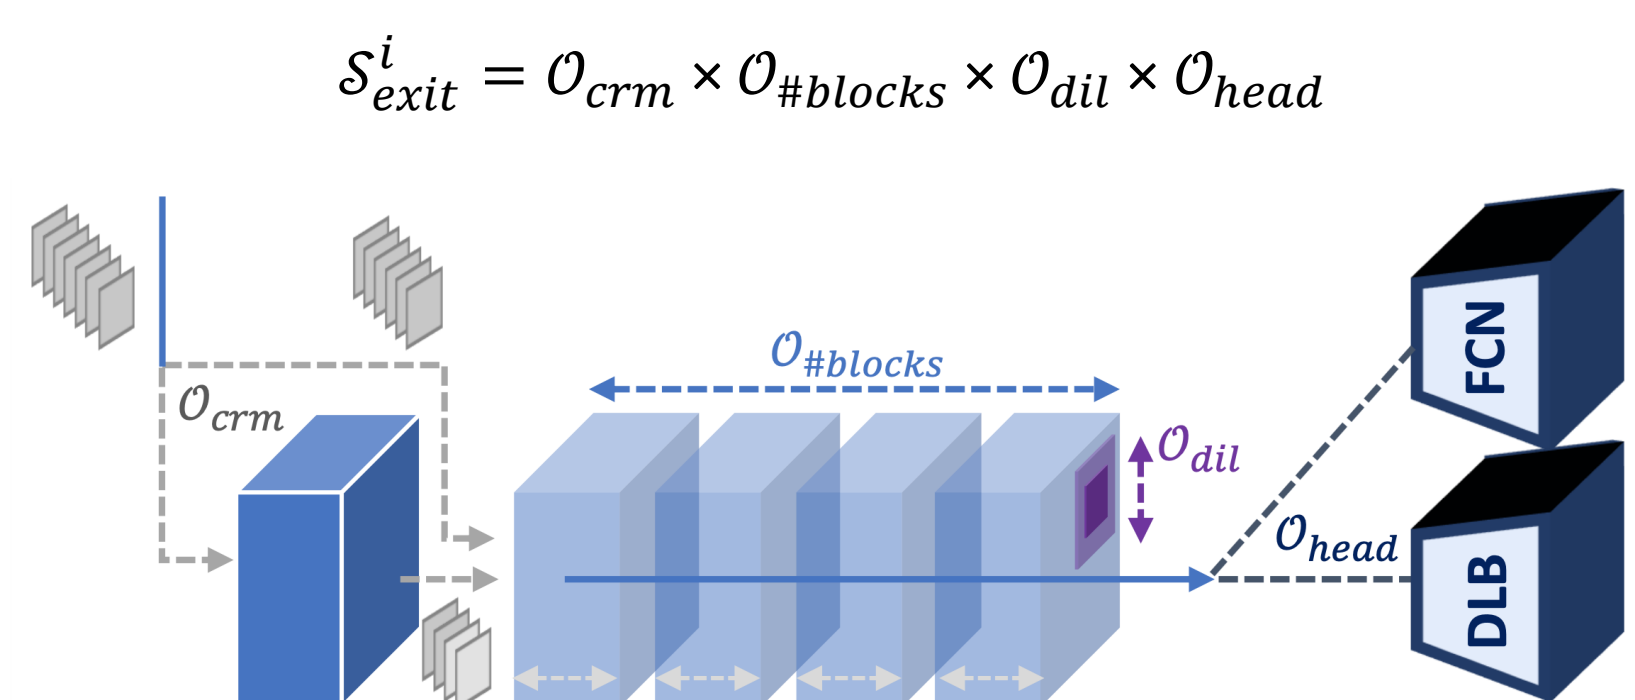
Given a (user-provided) semantic segmentation model, MESS framework measures the workload/latency of each layer to identify N candidate exit-points with approximately equidistant workload distribution.



Exit Architecture

Different exit-point (depths) benefit from different architectural configuration of exits. Shallow exits suffer from *limited receptive field* and *weak semantics*, whereas deeper ones can introduce significant computational overhead due to large incoming feature volume. Early-exits can incorporate:

1. **Channel Reduction module:** reducing the number of channels and this computational overhead.
2. **Extra Trainable Blocks:** To remedy weak semantics.
3. **Rapid Distillation Increase:** To increase receptive field.
4. **Segmentation head:** FCN and DeepLab-based segmentation output layers supported.



MESS Training

MESS networks are trained by a novel 2-stage scheme:

Stage1 - Exit-aware pre-training: trains the backbone in an exit-aware manner, pushing the extraction of semantically strong features to shallow layers, without committing to any exit configuration, or affecting the accuracy of deeper exits:

$$\mathcal{L}_{\text{pretrain}}^{\text{batch}(j)} = \sum_{i=1}^{N-1} \mathbb{1}(j \bmod i = 0) \cdot \mathcal{L}_{\text{CE}}(\mathbf{y}_i, \hat{\mathbf{y}}) + \mathcal{L}_{\text{CE}}(\mathbf{y}_N, \hat{\mathbf{y}})$$

Stage2 - Frozen backbone KD: Trains all candidate exit configurations, without any interference with the backbone (or between them). Allows massive parallelism during training, and easy interchange of exits upon deployment. Backed by a novel self-distillation scheme that filters the information propagating from the final exit to the earlier ones (allowing only “easy” pixels through).

$$\mathcal{L}_{\text{PFD}} = \sum_{i=1}^N \alpha \cdot \mathbb{1}(\hat{\mathbf{y}}_N = \hat{\mathbf{y}}) \mathcal{L}_{\text{CE}}(\mathbf{y}_i, \hat{\mathbf{y}}) + (1 - \alpha) \cdot \mathcal{L}_{\text{KL}}(\mathbf{y}_i, \mathbf{y}_N)$$

Search and Deployment

Upon deployment, MESS support different settings:

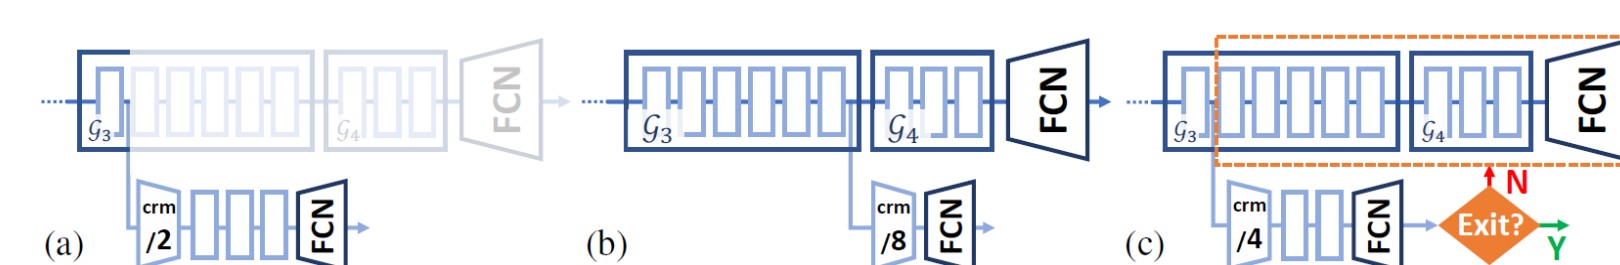
- (a) **Budgeted Inference:** static workload-lighter sub-models (up to an exit) are extracted and deployed.
- (b) **Anytime Inference:** every sample goes through several exits sequentially, providing an early actionable result while progressively refining its prediction.
- (c) **Input-Dependent Inference:** similar to anytime, but each sample dynamically finalises its output at a different depth, based on a difficulty-aware exit policy.

MESS instances can be optimised to:

- Minimise cost, given an accuracy constraint:

$$s^* = \arg \min_{s \in \mathcal{S}} \{\text{cost}(s) \mid \text{acc}(s) \geq th_{\text{acc}}\}$$
- Maximise accuracy, given a cost constraint:

$$s^* = \arg \max_{s \in \mathcal{S}} \{\text{acc}(s) \mid \text{cost}(s) \leq th_{\text{cost}}\}$$



Exit-Policy

We establish a mechanism to capture segmentation confidence on a *per-image* granularity, reducing the *per-pixel* confidence values provided with the model prediction. Our metric considers the percentage of pixels with high prediction confidence (above a tunable threshold) in the dense output (instead of naïve avg.).

$$c_i^{\text{img}} = \frac{1}{RC} \sum_{r=1}^R \sum_{c=1}^C \mathbb{1}(c_{r,c}^{\text{map}}(\mathbf{y}_i) \geq th_i^{\text{pix}})$$

Predictions in pixel along the semantic boundaries of each object are naturally under confident. Thus we introduce a mechanism to down-weight the contribution of these pixels to the final result.

$$\mathcal{M} = \text{erode}(\text{cannyEdge}(\hat{\mathbf{y}}_i), s_i)$$

$$\widehat{c}_{r,c}^{\text{map}}(\mathbf{y}_i) = \begin{cases} \text{median}(c_{w_r,w_c}^{\text{map}}(\mathbf{y}_i)) & \text{if } \mathcal{M}_{r,c} = 1 \\ c_{r,c}^{\text{map}}(\mathbf{y}_i) & \text{otherwise} \end{cases}$$

At runtime, a *per-exit tunable threshold* is used to evaluate whether each sample should exit-early or continuer to the next (deeper) exit: $c_i^{\text{img}} \geq th_i^{\text{img}}$

Results

MESS End-to-End Evaluation

Method	Backbone*	Head	Search Targets			Results: MS COCO			Results: PASCAL VOC		
			Error	GFLOPs	miou	GFLOPs	miou	GFLOPs	miou	GFLOPs	GFLOPs
DRN [68]	(i)	ResNet50	FCN	Baseline	59.02%	138.63	39.96ms	72.23%	138.63	39.93ms	72.23%
	(ii)	ResNet50	FCN	min	≤ 1x	64.35%	113.65	37.53ms	79.09%	113.65	37.59ms
	(iii)	ResNet50	FCN	≤ 0.1%	min	58.91%	41.17	17.92ms	72.16%	44.81	18.63ms
	(iv)	ResNet50	FCN	≤ 1%	min	58.12%	34.53	15.11ms	71.29%	38.51	16.80ms
DLBV3 [6]	(v)	ResNet50	DLB	Baseline	64.94%	163.86	59.05ms	80.32%	163.86	59.06ms	80.32%
	(vi)	ResNet50	DLB	min	≤ 1x	65.52%	124.10	43.29ms	82.32%	124.11	43.30ms
	(vii)	ResNet50	DLB	≤ 0.1%	min	64.86%	69.84	24.81ms	80.21%	65.29	24.14ms
	(viii)	ResNet50	DLB	≤ 1%	min	64.03%	57.01	20.83ms	79.30%	50.29	20.11ms
segMBNetV2 [50]	(ix)	MobileNetV2	FCN	Baseline	54.24%	8.78	67.04ms	69.68%	8.78	67.06ms	69.68%
	(x)	MobileNetV2	FCN	min	≤ 1x	57.49%	8.10	56.05ms	74.22%	8.10	56.09ms
	(xi)	MobileNetV2	FCN	≤ 0.1%	min	54.18%	4.05	40.97ms	69.61%	3.92	32.79ms
	(xii)	MobileNetV2	FCN	≤ 1%	min	53.24%	3.48	38.83ms	68.80%	3.60	31.40ms

*Dilated network [68] based on backbone CNN. *Measured on: GTX for ResNet50 and AGX for MobileNetV2 backbone.

MESS Training Evaluation

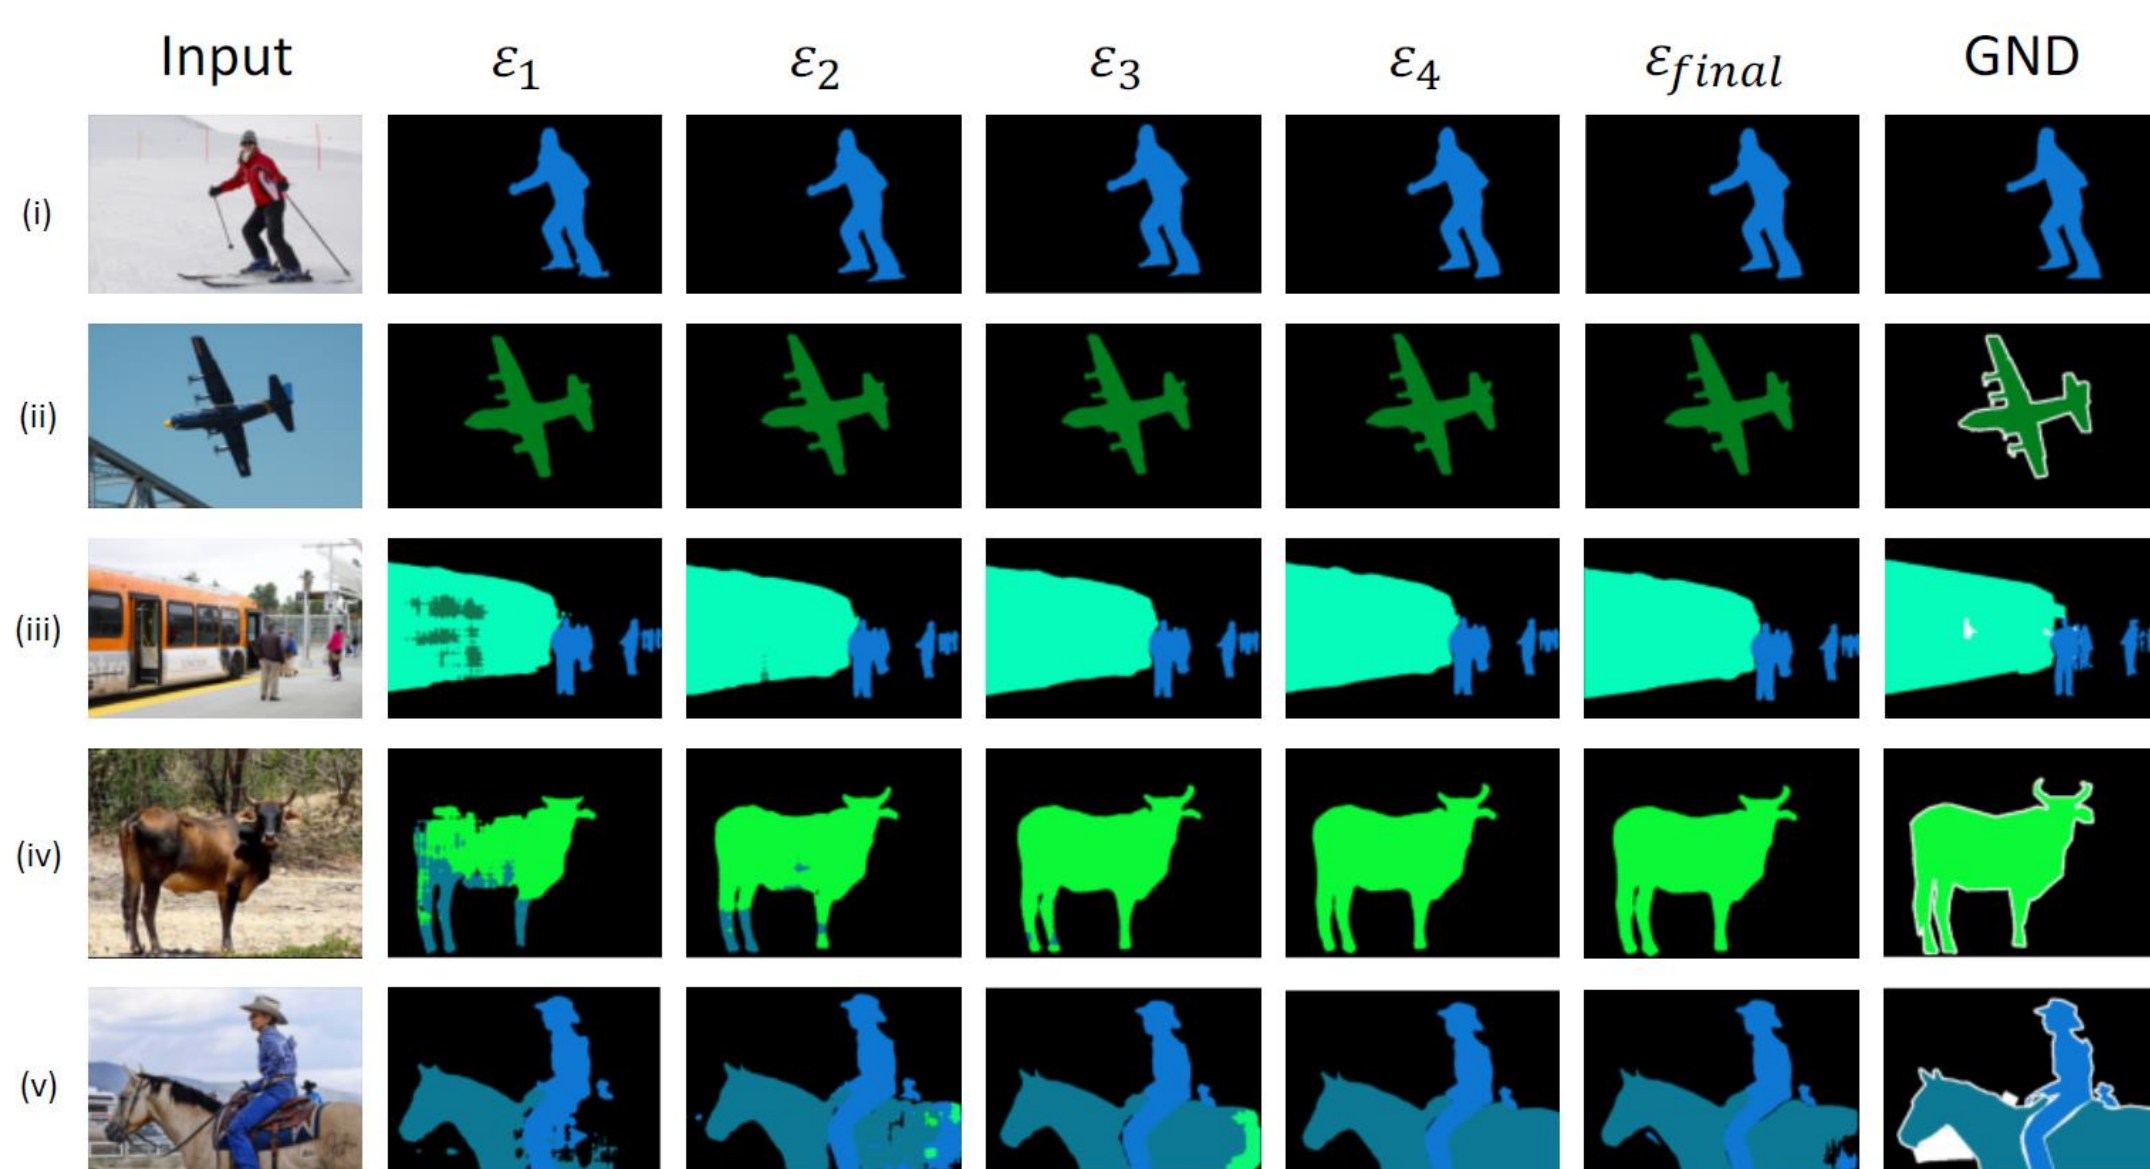
stage 1:	Method	Init.	Loss	Results: MS COCO					
				$\mathcal{L}_{\text{CE}}(\mathcal{E}_1)$	\mathcal{E}_1	\mathcal{E}_2	\mathcal{E}_3	\mathcal{E}_4	\mathcal{E}_5
(i) Baseline Init.	ImageNet	Baseline	59.02%	29.02%	40.67%	48.64%	51.69%	55.34%	58.35%
	E2E [21,53]	ImageNet	$\mathcal{L}_{\text{CE}}(\mathcal{E}_1) + \dots + \mathcal{L}_{\text{CE}}(\mathcal{E}_5)$	29.02%	40.67%	48.64%	51.69%	55.34%	58.35%
	Exit-aware Init.	ImageNet	$\mathcal{L}_{\text{CE}}(\mathcal{E}_1) + \dots + \mathcal{L}_{\text{CE}}(\mathcal{E}_5)$	28.21%	39.61%	47.13%	50.81%	56.11%	59.90%
	Our (25,28)	ImageNet	$\mathcal{L}_{\text{CE}}(\mathcal{E}_1) + \dots + \mathcal{L}_{\text{CE}}(\mathcal{E}_5)$	23.94%	31.50%	38.24%	44.73%	54.32%	59.02%
(ii) Frozen	Our (84,3)	ImageNet	$\mathcal{L}_{\text{CE}}(\mathcal{E}_1) + \dots + \mathcal{L}_{\text{CE}}(\mathcal{E}_5)$	32.40%	43.34%	50.81%	53.73%	57.9%	59.90%
	Our (84,3)	ImageNet	$\mathcal{L}_{\text{CE}}(\mathcal{E}_1) + \dots + \mathcal{L}_{\text{CE}}(\mathcal{E}_5)$	32.40%	43.34%	50.81%	53.73%	57.9%	59.90%

*Experiments repeated 3 times. The sample sizes in mean IoU is at most ± 0.09 in all cases.



stage 2:

Method	Loss	DRN-50			MobileNetV2		
		\mathcal{E}_1	\mathcal{E}_2	\mathcal{E}_3	\mathcal{E}_1	\mathcal{E}_2	\mathcal{E}_3
CE	CE	49.90%	55.40%	58.96%	31.56%	41.57%	51.59%
KD [19]	KD	50.33%	55.67%	59.08%	31.04%	41.93%	51.66%
SelfDistill [73,48,72]	CE+KD	50.66%	55.91%	58.84%	32.08%	41.96%	51.58%
Our (84,3)	PFD	51.02%	56.21%	59.46%	33.36%	42.95%	52.20%


CE=Cross-entropy, KD=Knowledge Distillation, PFD=Positive Filtering Distillation



Segmentation predictions from consecutive exits of a MESS instance based on DRN-50.



Comparison of distillation schemes used during MESS 2nd-stage training (frozen backbone).



Prediction confidence visualization (before and after) proposed refinement.